

## 02 Delta Lake

### Delta Lake

Read the original research paper for better understanding: <https://www.databricks.com/wp-content/uploads/2020/08/p975-armbrust.pdf>

### Introduction to Delta Lake

Delta Lake is an **open-source storage layer** that brings reliability to data lakes by enabling **ACID transactions**, **scalable metadata handling**, and unified **batch and streaming** data processing. It transforms traditional data lakes into robust **Lakehouse architectures**, combining the scalability of data lakes with the reliability of data warehouses.

- **Origin:** Developed by Databricks in 2016 and open-sourced in 2019 under the Linux Foundation. It's not controlled by any single company.
- **Core Purpose:** Addresses data lake challenges like lack of ACID compliance, schema enforcement, and time travel, without requiring data movement.
- **Key Benefits:**
  - **ACID Transactions:** Ensures atomicity, consistency, isolation, and durability for reliable updates.
  - **Schema Enforcement & Evolution:** Validates data on write and allows non-breaking schema changes.
  - **Time Travel:** Query historical data versions for audits, rollbacks, or ML reproducibility.
  - **Unified Processing:** Handles batch, streaming, and interactive queries seamlessly.
  - **Scalability:** Manages petabyte-scale tables with billions of files.

**Latest Version (as of September 2025):** Delta Lake 4.0.0, released on June 6, 2025, built on Apache Spark 4.0. This is the largest release to date, focusing on performance, reliability, and ease of use.

### Open Table Formats and Lakehouse Capabilities

Before diving into Delta Lake, understand **open table formats (OTFs)**: These add a **logical/metadata layer** on top of data lakes (e.g., stored in Parquet files on S3, ADLS, GCS) to enable Lakehouse features like SQL querying, transactions, and governance.

- **Why OTFs?:** Raw data lakes are cheap and scalable but lack performance, reliability, and SQL-like features. OTFs create a "virtual table" (schema + metadata) without moving data, allowing analysts to query as if it were a warehouse.
- **Popular OTFs:**
  - **Delta Lake:** Most mature for Spark-centric environments; backbone of Databricks and Microsoft Fabric.
  - **Apache Iceberg:** Strong in multi-engine support (e.g., Snowflake, AWS Athena); excels in schema evolution and large-scale analytics.
  - **Apache Hudi:** Optimized for real-time streaming, upserts, and CDC (change data capture).

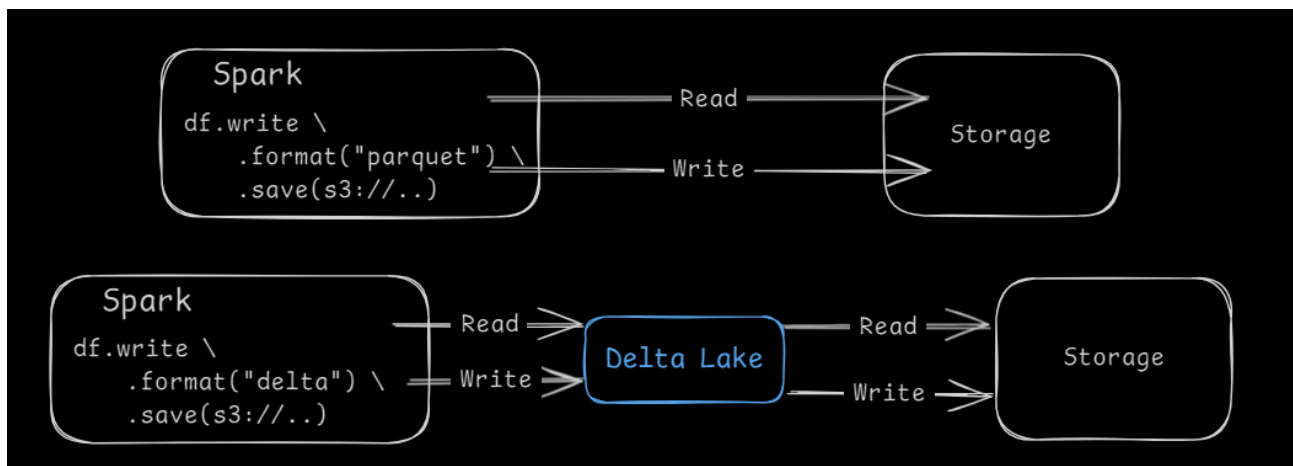
**Delta Lake's Position:** More mature than alternatives in Spark ecosystems, with tight integration. As of 2025, all three have ~80-90% feature parity, but choice depends on workload (e.g., Delta for unified

batch/streaming).

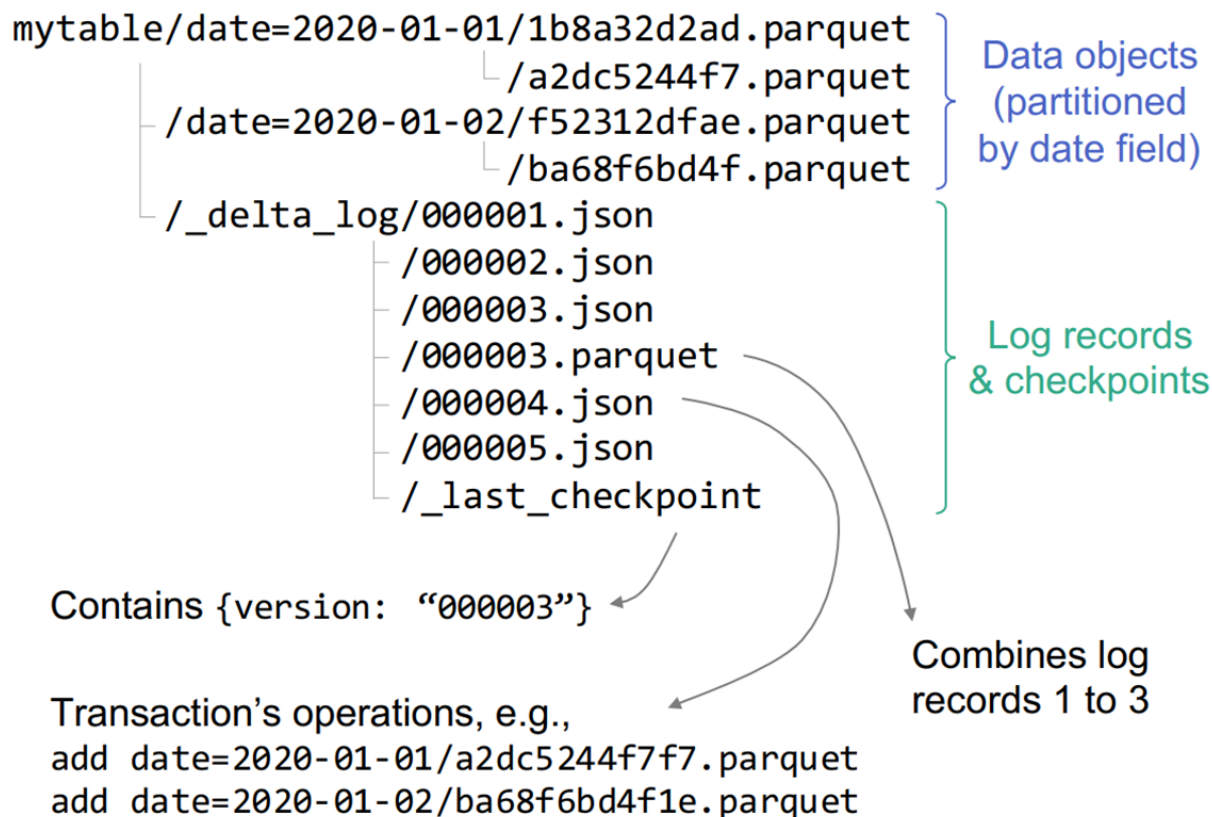
Comparison Table: Delta Lake vs. Iceberg vs. Hudi (2025)

Feature	Delta Lake	Apache Iceberg	Apache Hudi
ACID Transactions	Full support via transaction log	Full support via snapshots	Full via timeline-based commits
Time Travel	Yes (version/timestamp queries)	Yes (snapshot isolation)	Yes (timeline queries)
Schema Evolution	Enforcement + evolution (e.g., add/drop columns)	Advanced (partition evolution, type widening)	Good for upserts, but less flexible
Batch/Streaming	Unified (Structured Streaming)	Batch-focused; streaming via engines	Streaming-optimized (MoR for real-time)
Performance Opts	Z-Ordering, Liquid Clustering, Auto-Compaction	Hidden partitioning, Vectorized I/O	Indexing for upserts, MoR/COW modes
Multi-Engine Support	Spark, Flink, Trino, Hive, Snowflake; UniForm for Iceberg/Hudi	Broad (Spark, Trino, Flink, Snowflake, Athena)	Spark, Flink, Hive; good for CDC
Ecosystem	Databricks, Microsoft Fabric (default)	AWS, Snowflake, Google BigQuery	Uber-inspired; streaming-heavy
Maturity/Adoption	High (60%+ Fortune 500); Spark leader	Rapid growth (Netflix, AWS)	Strong in real-time (fintech, IoT)
Best For	Spark/Databricks users, mixed workloads	Multi-vendor lakehouses, analytics	Real-time ingestion, frequent updates

**Interoperability Note:** Delta Lake 4.0's **UniForm** allows Delta tables to be read by Iceberg/Hudi clients without data copying, reducing vendor lock-in.



### How Delta Lake Works



Delta Lake is **not a file format** (no `.delta` files)—it's a **protocol** built on **Parquet files** (columnar, compressed for analytics) + a **transactional log** (`_delta_log` folder).

- **Structure:**
  - **Parquet Files:** Store actual data (efficient for big data, like CSV/JSON but optimized).
  - **\_delta\_log Folder:** JSON files for transactions + Parquet checkpoints for fast reads. Uses **MVCC (Multi-Version Concurrency Control)** for optimistic concurrency.
    - Tracks: Metadata, file adds/deletes, schema changes, operations (insert/update/delete).

- Enables: ACID via atomic commits; prevents corruption in concurrent writes.
- **Conversion Process:**
  - Upgrade existing Parquet to Delta: Add the log folder (e.g., via `CONVERT TO DELTA` in Spark).
  - Example: Day 1 sales Parquet → Enable Delta → Log records metadata. Day 2 adds → Log appends changes.
- **Lakehouse Enablement:** The log turns raw lake data into a "managed table" with SQL features. Feels like a real database—no performance hit for users.

**Medallion Architecture:** Often used with Delta (Bronze: raw → Silver: cleaned → Gold: curated) for data quality in Databricks/Fabric.

### Key Features

- **ACID Transactions:** MVCC ensures serializability; supports **MERGE** for upserts (e.g., CDC).
- **Time Travel:** Query past versions: `SELECT * FROM table VERSION AS OF 5` or `TIMESTAMP AS OF '2025-01-01'`. Great for audits/rollbacks.
- **Optimizations** (via `OPTIMIZE` command):
  - **Auto-Compaction/Bin-Packing:** Merges small files (ideal size: 100-300 MB).
  - **Z-Ordering:** Clusters data by columns for faster queries (up to 70% speedup with Photon engine).
  - **Liquid Clustering** (Delta 3.2+): Adaptive, replaces partitioning; auto-tunes for dynamic data (no rewrites for key changes).
  - **Data Skipping/Caching:** Uses metadata stats to skip irrelevant files.
- **Schema Features:** Enforce on write; evolve (add columns, widen types like INT to LONG without rewrite in 4.0).
- **Security/Compliance:** Row-level security (RLS), PII masking; supports GDPR/HIPAA.
- **Streaming:** Exactly-once semantics with Spark Structured Streaming.
- **Delta 4.0 Highlights** (June 2025):
  - **Type Widening:** Evolve types (e.g., INT to LONG) without data rewrite.
  - **Row Tracking Backfill:** Enable on existing tables for row-level lineage (e.g., `_metadata.row_id`).
  - **Variant Data Type:** Flexible ingestion without schema; faster with Spark Variant encoding.
  - **Delta Connect:** Supports Spark Connect for client-server (remote Spark from anywhere).
  - **Zero-Copy Convert:** From Iceberg tables (Spark 3.5+).
  - **AI-Driven Opts:** Auto-tune Z-Orders based on queries (preview).

**UniForm (from Delta 3.0, enhanced in 4.0):** Generates Iceberg/Hudi metadata asynchronously.

Enable: `ALTER TABLE table SET TBLPROPERTIES ('delta.uniform.iceberg.enabled' = 'true')`. Read Delta as Iceberg/Hudi—no copies.

### Ecosystem and Adoption

- **Platforms:** Default in **Databricks** (Photon engine for 70% faster queries) and **Microsoft Fabric** (OneLake shortcuts for external querying; Direct Lake Mode in Power BI).
- **Engines:** Spark (native), Flink, Trino, Hive, PrestoDB, Snowflake, BigQuery, Athena, Redshift.
- **APIs:** Scala, Java, Python, Rust (delta-rs 1.0+ for community Rust impl).

- **Adoption:** 10M+ monthly downloads; used by 60%+ Fortune 500. Case: 40% cost reduction, 70% faster queries in Azure Databricks audits.
- **Community:** 190+ developers from 70+ orgs; governed by Linux Foundation. Backward compatible (newer reads older tables).

#### Use Cases:

- **Analytics/BI:** Medallion pipelines in Fabric/Databricks.
- **Real-Time:** Streaming with exactly-once ingestion.
- **ML:** Feature stores with time travel for reproducible experiments.
- **Compliance:** Time travel for audits (e.g., GDPR).

#### Practical Example (PySpark in Databricks/Spark)

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("DeltaExample").getOrCreate()

# Create Delta table from Parquet
data = spark.read.parquet("path/to/parquet")
data.write.format("delta").saveAsTable("my_delta_table")

# Time travel query
spark.sql("SELECT * FROM my_delta_table VERSION AS OF 1").show()

# Optimize
spark.sql("OPTIMIZE my_delta_table ZORDER BY (column_name)")

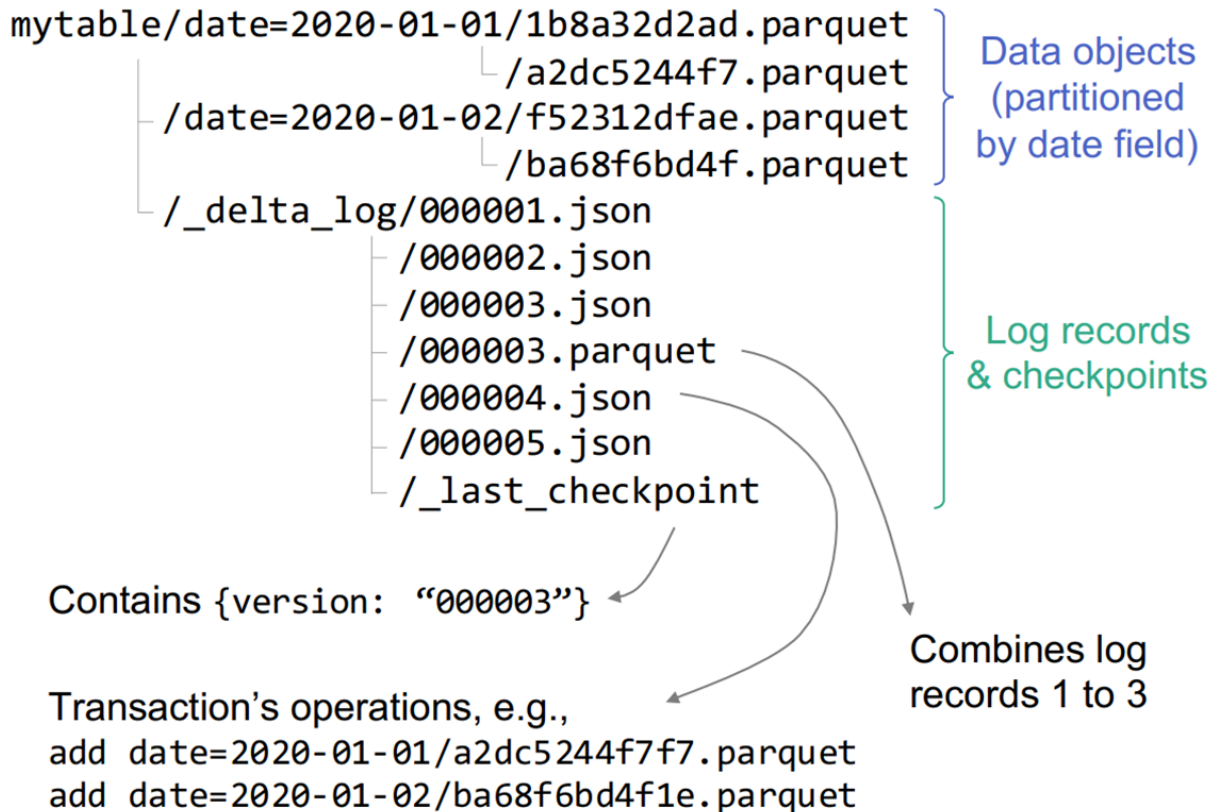
# Enable UniForm (Iceberg compat)
spark.sql("ALTER TABLE my_delta_table SET TBLPROPERTIES
('delta.uniform.iceberg.enabled' = 'true')")
```

#### Conclusion

Delta Lake is the **backbone of modern Lakehouses**, especially in Spark/Databricks ecosystems. With 4.0's UniForm and performance boosts, it's more interoperable than ever. For alternatives, evaluate based on engines (Iceberg for multi-vendor) or streaming needs (Hudi). Experiment in Databricks Community Edition for hands-on.

#### Further Reading:

- [Official Docs](#)
- [GitHub Releases](#)



## File Structure

### Data Files (Parquet)

- `mytable/date=2020-01-01/` and `date=2020-01-02/` folders contain the actual data
- Files like `1b8a32d2ad.parquet`, `a2dc5244f7.parquet` are the data files
- This shows **partitioning by date field** - data is organized by date for query optimization

### Delta Log Folder ( `_delta_log/` )

- Contains transaction history as JSON files: `000001.json`, `000002.json`, etc.
- `000003.parquet` - After many transactions, log entries get compacted into parquet format for efficiency
- `_last_checkpoint` - Points to the latest checkpoint for faster table state reconstruction

## How Transaction Logging Works

### Version Tracking

- Each JSON file represents a transaction/version
- `000003` means this is version 3 of the table
- The `_last_checkpoint` file contains `{"version": "000003"}` indicating the current state

**Transaction Operations** The diagram shows example operations stored in the log:

```
add date=2020-01-01/a2dc5244f7f7.parquet
add date=2020-01-02/ba68f6bd4f1e.parquet
```

## Key Concepts Illustrated

## Checkpoint Mechanism

- "Combines log records 1 to 3" means transactions 1, 2, and 3 have been consolidated
- This prevents the log from growing infinitely
- Makes table state reconstruction faster

## ACID Properties

- Every data file addition/removal is logged atomically
- You can reconstruct the exact table state at any version
- Enables time travel: `SELECT * FROM mytable VERSION AS OF 2`

## Why This Matters

This structure enables:

- **Concurrent reads/writes** without conflicts
- **Time travel queries** to any previous version
- **Fast metadata operations** (schema changes, partition info)
- **Data integrity** through complete audit trail

The genius is that while your data lives in simple parquet files, the `_delta_log` provides database-like ACID guarantees and versioning on top of cheap object storage.